

Deep Learning, Computer Vision, NLP & Reinforcement Learning

100-Question Practice MCQ Bank

No Answers · Self-Assessment Edition

University of Buea — COT | TCHAPGA TCHITO C. | 2025/2026

Course	CEC616 Deep Learning / CV / NLP / RL — 2025/2026
Topics	ANN · CNN · YOLO & OpenCV · Teachable Machine · MediaPipe · NLP (STT/TTS/Sentiment/Embeddings) · Q-Learning
Total Questions	100 MCQs across 7 sections
Instructions	Write A, B, C, or D on the Answer line under each question.

MY SCORE TRACKER						
ANN	CNN	CV	TM	MP	NLP	RL
___/15	___/15	___/15	___/8	___/7	___/20	___/20

SECTION A — Artificial Neural Networks (ANN)

Neurons, activation functions, backpropagation, optimizers, regularization, loss functions.

Q1. What is the basic computational unit of an ANN called?

[ANN — Fundamentals]

- A. Neuron / Perceptron
- B. Node Cluster
- C. Synapse Engine
- D. Weight Matrix

Answer: _____

Q2. Which function introduces non-linearity into a neural network?

[ANN — Activation Functions]

- A. Loss function
- B. Activation function
- C. Optimizer
- D. Regularizer

Answer: _____

Q3. What does the bias term in a neuron do?

[ANN — Fundamentals]

- A. Scales the input
- B. Shifts the activation function
- C. Normalizes output
- D. Removes noise

Answer: _____

Q4. Which algorithm is used to compute gradients for updating weights in an ANN?

[ANN — Training]

- A. Forward propagation
- B. Backpropagation
- C. K-means
- D. Gradient boosting

Answer: _____

Q5. The vanishing gradient problem is most commonly associated with which activation function?

[ANN — Training]

- A. ReLU
- B. Sigmoid
- C. Leaky ReLU
- D. Swish

Answer: _____

Q6. What is the purpose of the softmax activation function?

[ANN — Activation Functions]

- A. Binary classification
- B. Convert logits to a probability distribution over classes
- C. Regression output
- D. Feature scaling

Answer: _____

Q7. Which of the following is an optimizer that uses adaptive learning rates?

[ANN — Optimizers]

- A. SGD
- B. Adam
- C. Plain Batch Gradient Descent
- D. Momentum-only

Answer: _____

Q8. What is "overfitting" in the context of an ANN?

[ANN — Generalization]

- A. Model performs well on both train and test data
- B. Model performs well on training data but poorly on unseen data
- C. Model fails to converge
- D. Model has too few parameters

Answer: _____

Q9. Dropout is a technique used to:

[ANN — Regularization]

- A. Increase the number of neurons
- B. Prevent overfitting by randomly disabling neurons during training
- C. Speed up forward propagation
- D. Normalize inputs

Answer: _____

Q10. Which loss function is typically used for multi-class classification?

[ANN — Loss Functions]

- A. Mean Squared Error
- B. Categorical Cross-Entropy
- C. Hinge Loss
- D. Huber Loss

Answer: _____

Q11. What is an epoch in neural network training?

[ANN — Training]

- A. One forward pass only
- B. One complete pass through the entire training dataset
- C. One mini-batch update
- D. One layer of the network

Answer: _____

Q12. Which layer type connects every neuron to every neuron in the next layer?

[ANN — Architecture]

- A. Convolutional layer
- B. Pooling layer
- C. Fully connected (Dense) layer
- D. Recurrent layer

Answer: _____

Q13. Batch normalization helps primarily by:

[ANN — Training]

- A. Reducing internal covariate shift and stabilizing training
- B. Removing the need for activation functions

- C. Replacing backpropagation
- D. Increasing the dataset size

Answer: _____

Q14. What does the term "learning rate" control in gradient descent?

[ANN — Training]

- A. Number of layers
- B. Size of the steps taken towards minimizing the loss
- C. Number of epochs
- D. Batch size

Answer: _____

Q15. Which of these is NOT a common activation function?

[ANN — Fundamentals]

- A. ReLU
- B. Tanh
- C. Sigmoid
- D. Gradient Descent

Answer: _____

SECTION B — Convolutional Neural Networks (CNN)

Convolution, pooling, padding/stride, architectures (LeNet, ResNet, U-Net), transfer learning.

Q16. What is the primary purpose of a convolutional layer?

[CNN — Layers]

- A. Reduce dataset size
- B. Extract spatial features using learnable filters
- C. Normalize pixel values
- D. Classify the final output

Answer: _____

Q17. What does a pooling layer (e.g., max pooling) primarily achieve?

[CNN — Layers]

- A. Increases spatial dimensions
- B. Reduces spatial dimensions while retaining key features
- C. Adds more channels
- D. Normalizes the gradient

Answer: _____

Q18. In a CNN, what does "stride" refer to?

[CNN — Layers]

- A. The depth of the filter
- B. The number of pixels the filter moves at each step
- C. The size of the kernel
- D. The number of channels

Answer: _____

Q19. What is "padding" used for in CNNs?

[CNN — Layers]

- A. Reduces overfitting
- B. Preserves spatial dimensions of feature maps at the borders
- C. Normalizes activations
- D. Adds dropout

Answer: _____

Q20. Which famous CNN architecture introduced residual (skip) connections?

[CNN — Architectures]

- A. AlexNet
- B. VGGNet
- C. ResNet
- D. LeNet

Answer: _____

Q21. What is the role of the Flatten layer in a CNN?

[CNN — Layers]

- A. Converts a 2D feature map into a 1D vector for the dense layers
- B. Performs convolution
- C. Performs pooling
- D. Normalizes the data

Answer: _____

Q22. Which CNN architecture was specifically designed for handwritten digit recognition (MNIST)?

[CNN — Architectures]

- A. AlexNet
- B. LeNet-5
- C. VGG-16
- D. ResNet-50

Answer: _____

Q23. What is "receptive field" in a CNN?

[CNN — Theory]

- A. The output layer size

- B. The region of the input image a particular CNN neuron is influenced by
- C. The number of filters
- D. The batch size

Answer: _____

Q24. Data augmentation in CNN training is used to:

[CNN — Training]

- A. Reduce training time
- B. Artificially increase training data diversity to reduce overfitting
- C. Remove noisy labels
- D. Increase the number of layers

Answer: _____

Q25. Which of these is a common CNN-based image classification dataset?

[CNN — Datasets]

- A. CIFAR-10
- B. UCI Iris
- C. Boston Housing
- D. Titanic

Answer: _____

Q26. What does the term "feature map" refer to in a CNN?

[CNN — Theory]

- A. The raw input image
- B. The output of applying a filter (convolution) to the input
- C. The loss curve
- D. The confusion matrix

Answer: _____

Q27. Global Average Pooling (GAP) is often used to:

[CNN — Layers]

- A. Replace the Flatten + Dense layers, reducing parameters before classification
- B. Increase image resolution
- C. Apply data augmentation
- D. Compute the loss

Answer: _____

Q28. Which type of CNN architecture is commonly used for image segmentation?

[CNN — Architectures]

- A. U-Net
- B. AlexNet
- C. LeNet

D. Perceptron

Answer: _____

Q29. In transfer learning with CNNs, "freezing" layers means:

[CNN — Transfer Learning]

- A. Deleting layers from the model
- B. Preventing their weights from being updated during training
- C. Adding new layers
- D. Doubling the learning rate

Answer: _____

Q30. What is the typical input shape for an RGB image of size 224×224?

[CNN — Fundamentals]

- A. (224, 224)
- B. (224, 224, 1)
- C. (224, 224, 3)
- D. (3, 3, 224)

Answer: _____

SECTION C — Computer Vision: YOLO & OpenCV

YOLO object detection, IoU, NMS, anchor boxes, mAP, OpenCV functions (imread, cvtColor, Canny, Haar Cascades).

Q31. What does YOLO stand for?

[CV — YOLO]

- A. You Only Live Once
- B. You Only Look Once
- C. Your Object Locator Optimized
- D. Yield Output Layer Operation

Answer: _____

Q32. YOLO performs object detection by:

[CV — YOLO]

- A. Running a sliding window classifier multiple times
- B. Dividing the image into a grid and predicting bounding boxes and class probabilities in a single forward pass
- C. Using only edge detection
- D. Segmenting pixels individually

Answer: _____

Q33. In object detection, what does "IoU" (Intersection over Union) measure?

[CV — Object Detection]

- A. Training speed
- B. Overlap between predicted and ground-truth bounding boxes
- C. Number of classes
- D. Model size

Answer: _____

Q34. Non-Maximum Suppression (NMS) in object detection is used to:

[CV — Object Detection]

- A. Increase the number of bounding boxes
- B. Remove duplicate/overlapping bounding boxes, keeping the most confident one
- C. Normalize the image
- D. Resize the input

Answer: _____

Q35. Which Python library is the standard tool for classical computer vision operations like edge detection and image filtering?

[CV — OpenCV]

- A. NumPy
- B. OpenCV
- C. Matplotlib
- D. Pandas

Answer: _____

Q36. In OpenCV, which function is used to read an image from disk?

[CV — OpenCV]

- A. cv2.load()
- B. cv2.imread()
- C. cv2.readImage()
- D. cv2.open()

Answer: _____

Q37. What color format does OpenCV use by default when reading images?

[CV — OpenCV]

- A. RGB
- B. BGR
- C. HSV
- D. Grayscale

Answer: _____

Q38. Which OpenCV function converts an image to grayscale?

[CV — OpenCV]

- A. cv2.gray()

- B. `cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)`
- C. `cv2.toGray()`
- D. `cv2.grayscale()`

Answer: _____

Q39. The Canny algorithm in OpenCV is used for:

[CV — OpenCV]

- A. Face detection
- B. Edge detection
- C. Object tracking
- D. Color segmentation

Answer: _____

Q40. What is the purpose of Haar Cascade classifiers in OpenCV?

[CV — OpenCV]

- A. Image compression
- B. Real-time object detection (e.g., face detection) using pre-trained feature-based classifiers
- C. Color correction
- D. Video encoding

Answer: _____

Q41. In YOLOv8, which library/framework is commonly used to load and run the model?

[CV — YOLO]

- A. Ultralytics
- B. OpenAI Gym
- C. NLTK
- D. Scrapy

Answer: _____

Q42. What does "anchor box" refer to in object detection models?

[CV — Object Detection]

- A. A fixed reference bounding box of predefined aspect ratio used to predict object location
- B. The center pixel of an image
- C. A type of pooling layer
- D. A loss function

Answer: _____

Q43. Which metric is commonly used to evaluate object detection model accuracy across all classes?

[CV — Object Detection]

- A. Accuracy
- B. mAP (mean Average Precision)
- C. R-squared

D. Silhouette Score

Answer: _____

Q44. What is the role of the confidence threshold in YOLO inference?

[CV — YOLO]

- A. Filters out predictions below a certain probability score
- B. Sets the learning rate
- C. Defines image resolution
- D. Controls batch size

Answer: _____

Q45. Which OpenCV function is used to draw a rectangle (e.g., a bounding box) on an image?

[CV — OpenCV]

- A. cv2.rectangle()
- B. cv2.box()
- C. cv2.drawRect()
- D. cv2.bound()

Answer: _____

SECTION D — Teachable Machine

Google's no-code ML training tool — Image, Sound, Pose projects, TensorFlow.js export.

Q46. Teachable Machine, developed by Google, allows users to train ML models primarily through:

[Teachable Machine]

- A. Writing TensorFlow code
- B. A no-code, browser-based drag-and-drop interface
- C. Command-line scripting only
- D. SQL queries

Answer: _____

Q47. Which types of models can Teachable Machine train?

[Teachable Machine]

- A. Image, Sound, and Pose models
- B. Only image models
- C. Only NLP models
- D. Only regression models

Answer: _____

Q48. What underlying technology does Teachable Machine use to run models directly in the browser?

[Teachable Machine]

- A. PyTorch Mobile
- B. TensorFlow.js
- C. ONNX Runtime
- D. Core ML

Answer: _____

Q49. In Teachable Machine, what are "classes" in an image project?

[Teachable Machine]

- A. CSS classes for styling
- B. Categories/labels the model learns to distinguish between
- C. Python classes
- D. Hardware specifications

Answer: _____

Q50. Which export format allows a Teachable Machine model to run on a website?

[Teachable Machine]

- A. .csv
- B. TensorFlow.js (JSON + weights)
- C. .docx
- D. .pdf

Answer: _____

Q51. What is a key advantage of Teachable Machine for beginners?

[Teachable Machine]

- A. Requires advanced calculus knowledge
- B. No coding or ML expertise needed to train a working classifier
- C. Only works with massive datasets
- D. Requires a GPU cluster

Answer: _____

Q52. In the Pose project type of Teachable Machine, what is the model primarily detecting?

[Teachable Machine]

- A. Sound frequency
- B. Human body keypoints/posture
- C. Text sentiment
- D. Color histograms

Answer: _____

Q53. What is recommended for improving accuracy of a Teachable Machine image model?

[Teachable Machine]

- A. Using only one sample per class
- B. Capturing diverse samples per class (different angles, lighting, backgrounds)

- C. Disabling data collection
- D. Reducing the number of classes to one

Answer: _____

SECTION E — MediaPipe

Real-time perception pipelines — Hands, Face Mesh, Pose, Holistic, graph-based calculators.

Q54. MediaPipe, developed by Google, is primarily used for:

[MediaPipe]

- A. Database management
- B. Real-time perception pipelines (e.g., hand, face, pose tracking)
- C. Natural language translation only
- D. Cloud storage

Answer: _____

Q55. Which MediaPipe solution is commonly used to track 21 hand landmarks in real time?

[MediaPipe]

- A. MediaPipe Hands
- B. MediaPipe Face Mesh
- C. MediaPipe Pose
- D. MediaPipe Holistic

Answer: _____

Q56. MediaPipe Face Mesh estimates approximately how many 3D facial landmarks?

[MediaPipe]

- A. 10
- B. 68
- C. 468
- D. 1000

Answer: _____

Q57. What is a key benefit of MediaPipe pipelines for real-time applications?

[MediaPipe]

- A. They require cloud GPU servers only
- B. They are optimized to run efficiently on-device (mobile/edge) with low latency
- C. They only work offline with no camera input
- D. They cannot process video, only static images

Answer: _____

Q58. MediaPipe Pose is primarily used to detect:

[MediaPipe]

- A. Text in images
- B. Full-body human pose landmarks (33 keypoints)
- C. Audio frequency
- D. Color palettes

Answer: _____

Q59. Which MediaPipe solution combines face, hand, and pose tracking into a single pipeline?

[MediaPipe]

- A. MediaPipe Holistic
- B. MediaPipe Objectron
- C. MediaPipe Selfie Segmentation
- D. MediaPipe Iris

Answer: _____

Q60. MediaPipe's underlying architecture is often described as a:

[MediaPipe]

- A. Single monolithic neural network
- B. Graph-based pipeline of modular, configurable calculators
- C. SQL-based pipeline
- D. Spreadsheet-based system

Answer: _____

SECTION F — NLP: Speech-to-Text, Text-to-Speech, Sentiment & Embeddings

STT/TTS systems, sentiment analysis (VADER), word embeddings (Word2Vec, BERT), tokenization.

Q61. Speech-to-Text (STT) systems convert:

[NLP — Speech-to-Text]

- A. Text into speech audio
- B. Spoken audio into written text
- C. Images into text
- D. Text into another language

Answer: _____

Q62. Which deep learning architecture is commonly used as the backbone for modern STT systems (e.g., Whisper)?

[NLP — Speech-to-Text]

- A. Decision Tree
- B. Transformer / Encoder-Decoder
- C. K-Means

D. Random Forest

Answer: _____

Q63. Text-to-Speech (TTS) systems are evaluated using which type of subjective quality metric?

[NLP — Text-to-Speech]

- A. Accuracy
- B. Mean Opinion Score (MOS)
- C. F1-score
- D. R-squared

Answer: _____

Q64. Which Python library/API is widely used for offline speech recognition prototyping?

[NLP — Speech-to-Text]

- A. SpeechRecognition
- B. NumPy
- C. OpenCV
- D. Pandas

Answer: _____

Q65. What does "sentiment analysis" determine from a piece of text?

[NLP — Sentiment Analysis]

- A. Grammar correctness
- B. The emotional tone/polarity (positive, negative, neutral)
- C. The text's language
- D. The word count

Answer: _____

Q66. Which is a classic lexicon-based approach to sentiment analysis?

[NLP — Sentiment Analysis]

- A. BERT
- B. VADER (Valence Aware Dictionary for sEntiment Reasoning)
- C. YOLO
- D. ResNet

Answer: _____

Q67. In NLP, what is a "word embedding"?

[NLP — Embeddings]

- A. A bag-of-words frequency count
- B. A dense vector representation of a word capturing semantic meaning
- C. A one-hot encoded sparse vector
- D. A grammatical parse tree

Answer: _____

Q68. Which embedding model introduced predicting a word from its context (or vice versa) using shallow neural networks?

[NLP — Embeddings]

- A. Word2Vec
- B. YOLO
- C. ResNet
- D. K-Means

Answer: _____

Q69. What does cosine similarity measure when comparing two word embeddings?

[NLP — Embeddings]

- A. The Euclidean distance only
- B. The angle/orientation similarity between two vectors
- C. The word frequency
- D. The sentence length

Answer: _____

Q70. Which embedding technique generates context-dependent (dynamic) embeddings, unlike Word2Vec's static embeddings?

[NLP — Embeddings]

- A. GloVe
- B. BERT
- C. TF-IDF
- D. Bag of Words

Answer: _____

Q71. What is "tokenization" in NLP preprocessing?

[NLP — Preprocessing]

- A. Encrypting the text
- B. Splitting text into smaller units such as words or subwords
- C. Converting text to speech
- D. Translating text

Answer: _____

Q72. Which library is widely used in Python for pre-trained Transformer-based NLP models (e.g., BERT, GPT)?

[NLP — Tools]

- A. Hugging Face Transformers
- B. OpenCV
- C. Scikit-image
- D. Matplotlib

Answer: _____

Q73. In sentiment analysis, a confusion matrix with classes Positive/Negative would NOT include which of the following metrics?

[NLP — Sentiment Analysis]

- A. Precision
- B. Recall
- C. mAP
- D. F1-score

Answer: _____

Q74. What is "TF-IDF" used for in text analysis?

[NLP — Embeddings]

- A. Measuring word importance in a document relative to a corpus
- B. Performing speech synthesis
- C. Detecting objects in images
- D. Tracking hand landmarks

Answer: _____

Q75. Which neural network type was historically used for sequential text data before Transformers became dominant?

[NLP — Architectures]

- A. CNN only
- B. RNN / LSTM
- C. K-Means
- D. Decision Tree

Answer: _____

Q76. What does "Speech-to-Text latency" typically refer to?

[NLP — Speech-to-Text]

- A. The model's file size
- B. The time delay between speaking and receiving the transcribed text
- C. The number of training epochs
- D. The vocabulary size

Answer: _____

Q77. In embedding spaces, the famous analogy "king – man + woman ≈ queen" demonstrates:

[NLP — Embeddings]

- A. Embeddings can be randomly generated
- B. Vector arithmetic captures semantic relationships between words
- C. Embeddings only work for nouns
- D. TTS quality

Answer: _____

Q78. Which of the following is an example of a cloud-based Text-to-Speech API?

[NLP — Text-to-Speech]

- A. Google Cloud Text-to-Speech
- B. OpenCV DNN
- C. YOLOv8
- D. Scikit-learn

Answer: _____

Q79. What is the main challenge that subword tokenization (e.g., Byte-Pair Encoding) solves?

[NLP — Preprocessing]

- A. Image resizing
- B. Handling out-of-vocabulary (OOV) and rare words efficiently
- C. Object detection accuracy
- D. Pose estimation

Answer: _____

Q80. Sentiment analysis applied to product reviews is an example of which broader NLP task category?

[NLP — Sentiment Analysis]

- A. Text classification
- B. Speech recognition
- C. Image captioning
- D. Object detection

Answer: _____

SECTION G — Reinforcement Learning: Q-Learning

Agent/environment/reward/policy, Bellman equation, epsilon-greedy, Deep Q-Learning (DQN).

Q81. In Reinforcement Learning, the entity that learns and makes decisions is called the:

[RL — Fundamentals]

- A. Environment
- B. Agent
- C. Reward
- D. Policy Network only

Answer: _____

Q82. What is a "reward" in Reinforcement Learning?

[RL — Fundamentals]

- A. A label for supervised training
- B. A scalar feedback signal indicating how good an action was

- C. The dataset size
- D. The number of episodes

Answer: _____

Q83. What does the "policy" in RL define?

[RL — Fundamentals]

- A. The reward function
- B. The mapping from states to actions that the agent follows
- C. The environment's transition probabilities
- D. The discount factor

Answer: _____

Q84. Q-Learning is classified as which type of Reinforcement Learning algorithm?

[RL — Q-Learning]

- A. Model-based, on-policy
- B. Model-free, off-policy
- C. Supervised learning
- D. Unsupervised clustering

Answer: _____

Q85. In Q-Learning, what does the "Q" in Q-table/Q-value represent?

[RL — Q-Learning]

- A. Quantity
- B. Quality of an action taken in a given state (expected future reward)
- C. Query
- D. Queue

Answer: _____

Q86. What is the Q-Learning update rule primarily based on?

[RL — Q-Learning]

- A. Backpropagation
- B. The Bellman equation
- C. Gradient boosting
- D. K-means clustering

Answer: _____

Q87. In the Q-Learning update equation, what does the discount factor (γ) control?

[RL — Q-Learning]

- A. The learning rate
- B. The importance given to future rewards relative to immediate rewards
- C. The exploration rate
- D. The number of states

Answer: _____

Q88. What does the exploration-exploitation trade-off in RL refer to?

[RL — Exploration]

- A. Choosing between trying new actions (explore) vs. choosing the best known action (exploit)
- B. Choosing between CNN and RNN
- C. Choosing the dataset size
- D. Choosing the batch size

Answer: _____

Q89. Which strategy is commonly used to balance exploration and exploitation in Q-Learning?

[RL — Exploration]

- A. Epsilon-greedy strategy
- B. Gradient descent
- C. Backpropagation
- D. Cross-validation

Answer: _____

Q90. What does the learning rate (α) control in the Q-Learning update rule?

[RL — Q-Learning]

- A. How much new information overrides old Q-values
- B. The discount applied to future rewards
- C. The number of episodes
- D. The size of the state space

Answer: _____

Q91. In RL terminology, what is a "state"?

[RL — Fundamentals]

- A. The agent's current situation/configuration in the environment
- B. The reward value
- C. The action taken
- D. The policy network's weights

Answer: _____

Q92. What is an "episode" in Reinforcement Learning?

[RL — Fundamentals]

- A. A single action
- B. A complete sequence of states, actions, and rewards from start to a terminal state
- C. One training epoch in supervised learning
- D. A single Q-table update

Answer: _____

Q93. Which of these best describes Deep Q-Learning (DQN)?

[RL — Deep RL]

- A. Q-Learning combined with a neural network to approximate Q-values for large/continuous state spaces
- B. Q-Learning without rewards
- C. Q-Learning using only a lookup table for all problems
- D. A purely supervised technique

Answer: _____

Q94. In Q-Learning, when the Q-table becomes impractically large (e.g., continuous state spaces), what is the typical solution?

[RL — Deep RL]

- A. Use a smaller reward
- B. Approximate the Q-function with a neural network (Deep Q-Network)
- C. Ignore some states entirely
- D. Remove the discount factor

Answer: _____

Q95. What does "off-policy" mean in the context of Q-Learning?

[RL — Q-Learning]

- A. The agent learns the value of the optimal policy independently of the actions it actually takes
- B. The agent only learns from human demonstrations
- C. The agent has no policy
- D. The agent cannot explore

Answer: _____

Q96. Which classic RL environment toolkit is commonly used to test Q-Learning algorithms (e.g., CartPole, FrozenLake)?

[RL — Tools]

- A. OpenAI Gym / Gymnasium
- B. OpenCV
- C. Hugging Face Transformers
- D. MediaPipe

Answer: _____

Q97. What typically happens to the exploration rate (epsilon) over training time in epsilon-greedy Q-Learning?

[RL — Exploration]

- A. It increases over time
- B. It decays/decreases over time, shifting from exploration to exploitation
- C. It stays constant always
- D. It becomes negative

Answer: _____

Q98. In the Bellman equation for Q-Learning, what does $\max_a Q(s', a)$ represent?

[RL — Q-Learning]

- A. The current reward
- B. The maximum expected future Q-value achievable from the next state
- C. The learning rate
- D. The initial Q-value

Answer: _____

Q99. A well-known real-world application of Q-Learning and Deep RL is:

[RL — Applications]

- A. Game-playing agents (e.g., Atari games, AlphaGo elements)
- B. Image classification only
- C. Static image segmentation
- D. Sentiment lexicon creation

Answer: _____

Q100. What is the terminal state in an RL episode?

[RL — Fundamentals]

- A. The very first state of the episode
- B. A state where the episode ends (e.g., goal reached or failure)
- C. The state with the highest reward always
- D. A randomly chosen state

Answer: _____
